

# Bayesian Statistics — Assignment 2

## On the Bayesian Lasso

Caio Peixoto

July 4, 2024

### 1 LASSO regression and the Laplace prior

The<sup>1</sup>

- Introduce notation for linear regression;
- Explain the Lasso regularization;
- Explain why the L1 norm produces shrinkage;
- Show how the LASSO estimator is the MAP estimator when using the Laplace prior;
- Explain the conditional structure of the Laplace prior;
- Comment on possible advantages of having a full probability distribution for the parameters;
- Talk about possible disadvantages of using the Laplace prior, allude to excessive shrinkage in the experiments;

### 2 The Gibbs sampler

The<sup>2</sup>

- Explain the sampling hierarchy suggested by Park and Casella.
- Modify the hierarchy to facilitate MAP estimation;
- Discuss how one should include information about the measurement noise parameter?
- Discuss why marginalize  $\mu$  in the computation and when that is and is not desirable;
- Show how to recover inferences about  $\mu$  in the marginalized case.

---

<sup>1</sup>Corresponds to items a) and b).

<sup>2</sup>Corresponds to items c) and d).

### 3 But can it actually shrink?

For<sup>3</sup> each scenario:

- Show table with Stan summary;
- Show trace plots of chains for all parameters;
- Show prior and posterior predictive distributions;
- Show graph with 95% confidence intervals for all parameters;

### 4 Choosing $\lambda$

The<sup>4</sup>

### 5 The “Huberised” LASSO

The<sup>5</sup>

---

<sup>3</sup>Corresponds to item e).

<sup>4</sup>Corresponds to item f).

<sup>5</sup>Corresponds to item g).